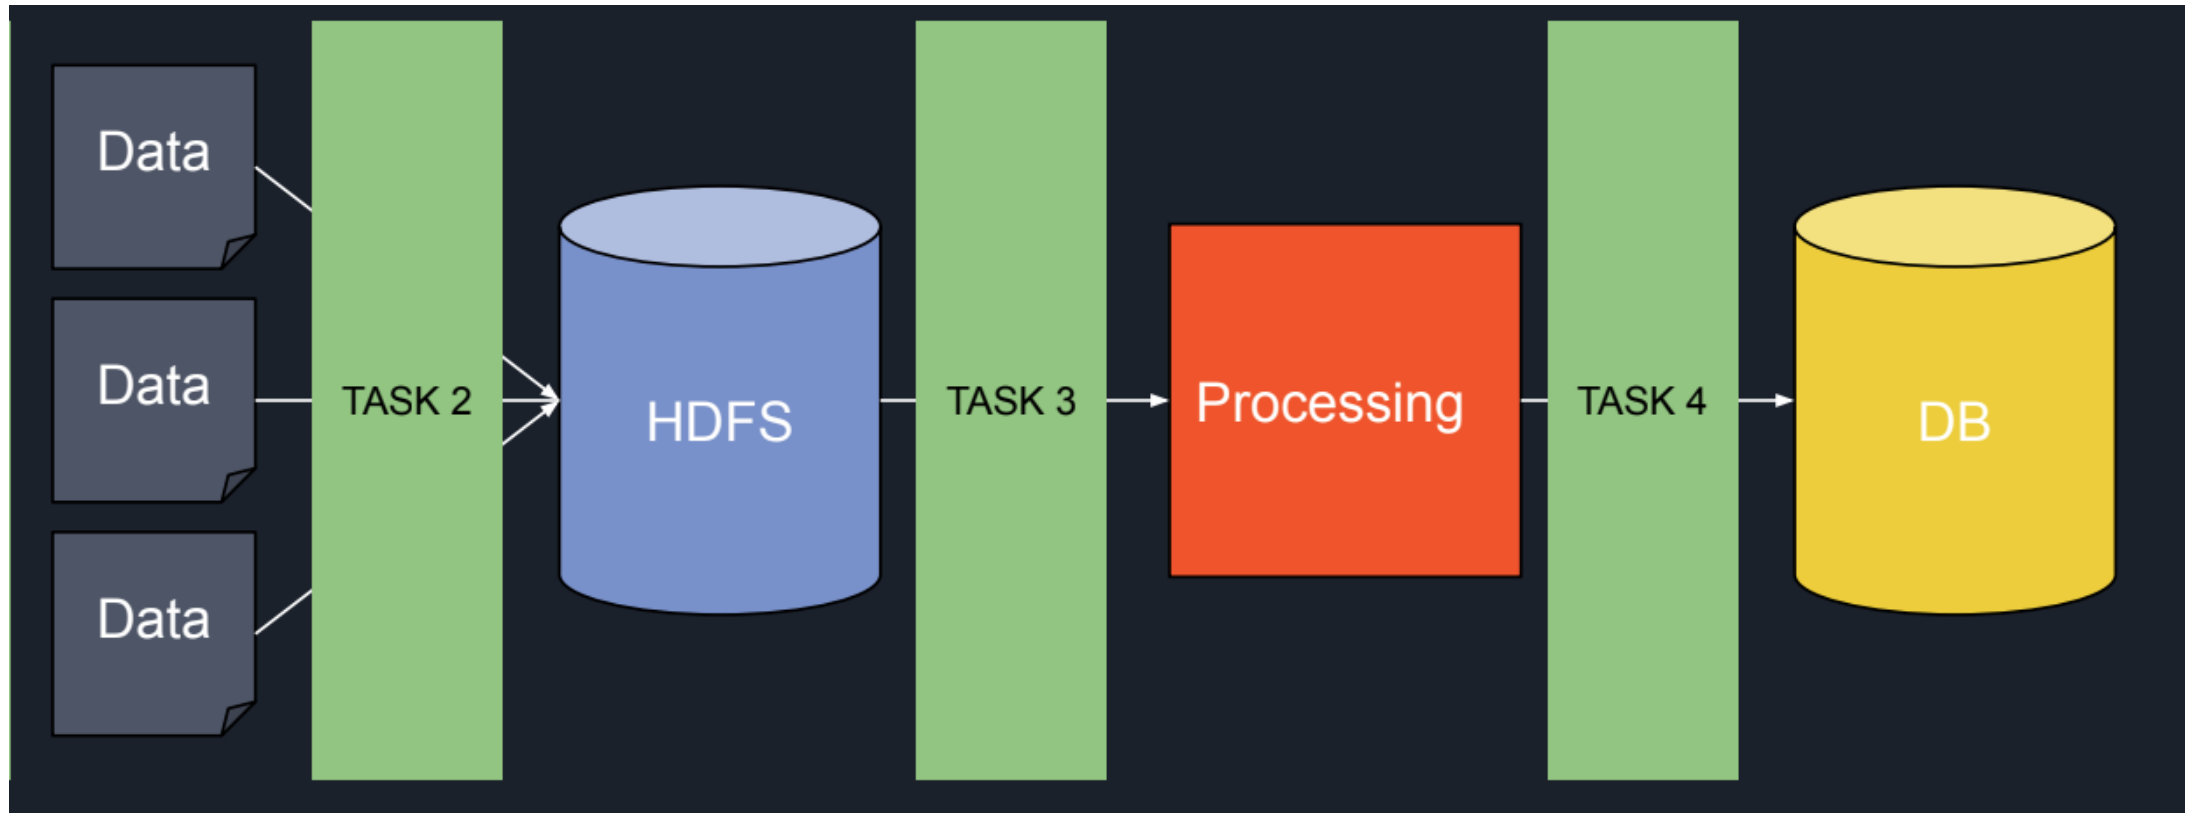


Apache Airflow



Content

- Introduction to the Airflow
- Core components of Airflow. (Web server, scheduler, meta store, executor*, Worker)
- DAG, Operators, Task/Task Instance, Workflow
- How Airflow works ?
- Airflow Useful Features
 - Subdags,
 - Taskgroups,
 - XComs,
 - Connections,
 - Variables,
 - Trigger Rules
 - Task flow API
- Working Demo, Troubleshooting

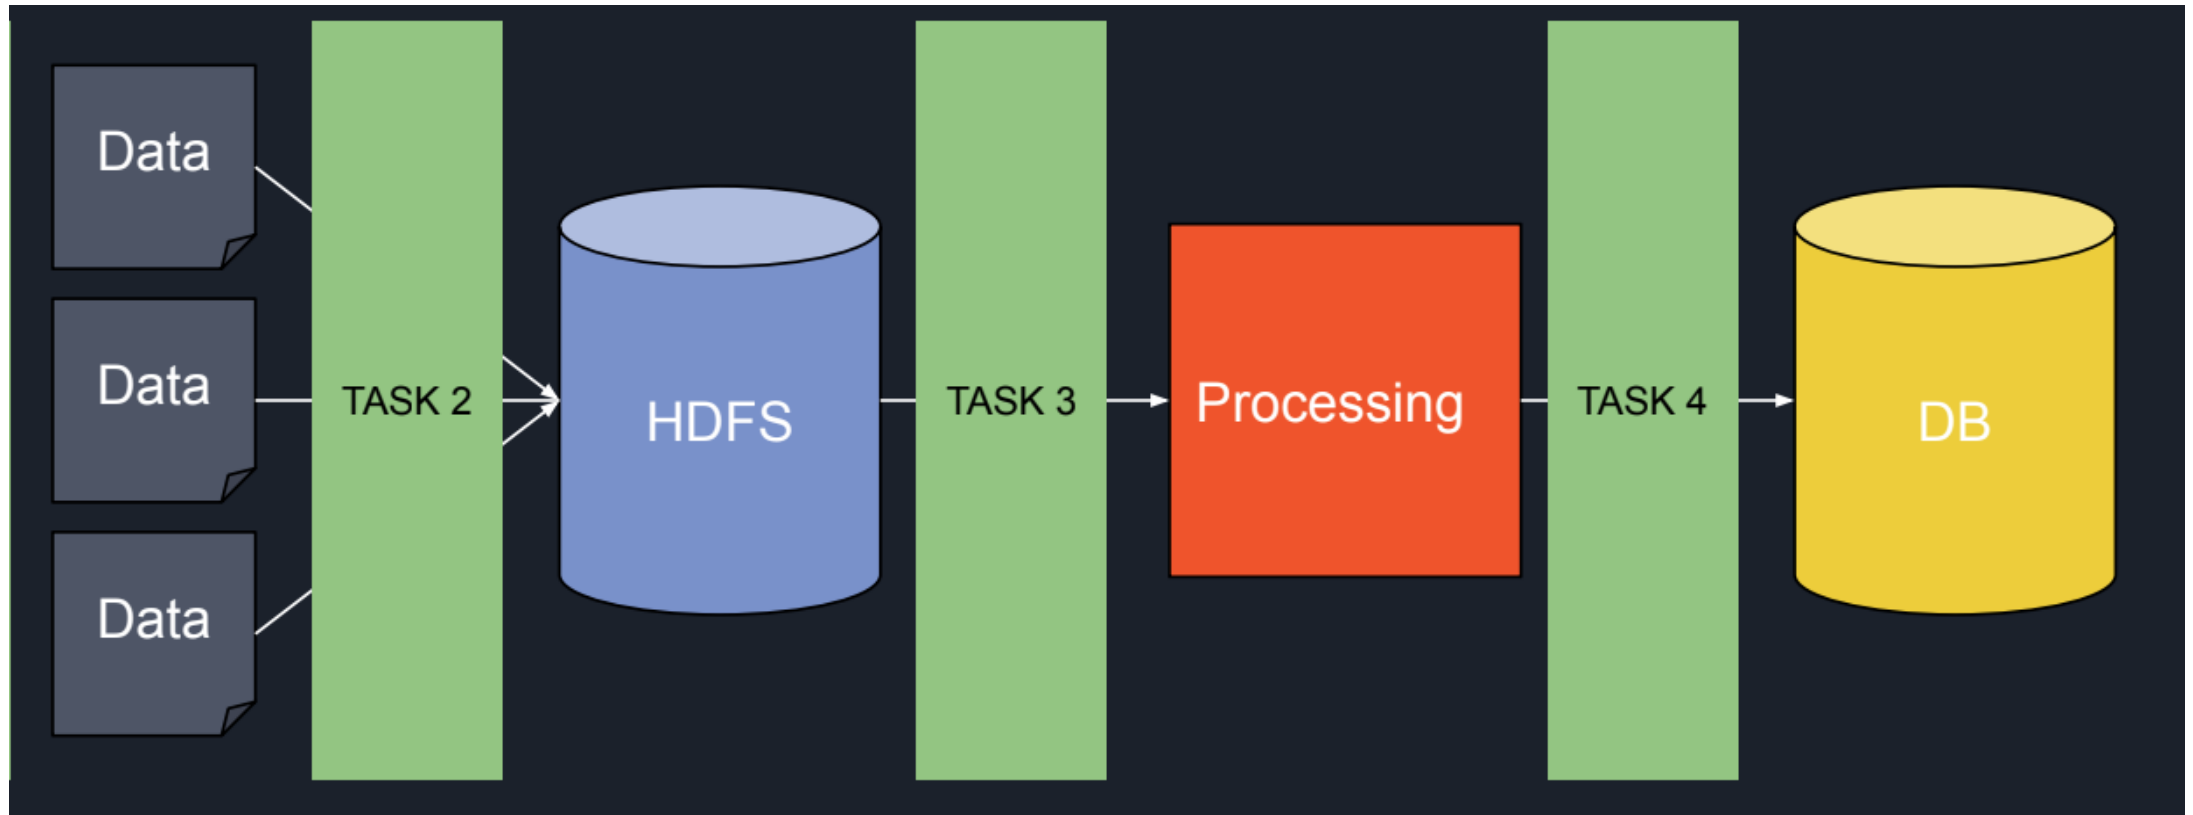


Let's imagine...

You are working for a company and you have a data processing pipeline to run every day at 9 AM which does the following:

1. Wait for files to come in a specific directory
 2. Store all the files into HDFS
 3. Run a Spark job to process those files
 4. Check the result of the job from the PostgreSQL database
- Problems?

Why Airflow?



Why Airflow?

Problems

- What if ... a file does not arrived in time?
- What if ... my Spark job failed?
- What if ... I have 1000 pipelines to execute?

Why Airflow

Airflow handles those problems and more

- Cron Replacement
- Fault Tolerant
- Dependency Rules
- Python Code
- Handle Task Failures
- Report / Alert on failures
- Extensible and modifiable
- Beautiful UI

Airflow is a perfect tool in order to create, monitor and manage your data pipelines.





What is Airflow




- **Apache Airflow is an open source platform to author, schedule and monitor workflows**



What Airflow is Not?

Airflow is not a data streaming solution

- Airflow is not in the scope of Apache Spark or Storm.
 - Primarily built to perform scheduled batch jobs
- 



Cloud providers -



Core Components of Airflow



Web Server



Scheduler



Metastore database



Executor



Worker

Key Concepts



DAG



Operator



Task




Task Instance



Workflow

- DAGS



operator

```
file = open("myfile", "r")  
print(f.read())
```

Click to add text

Operator


3 Types

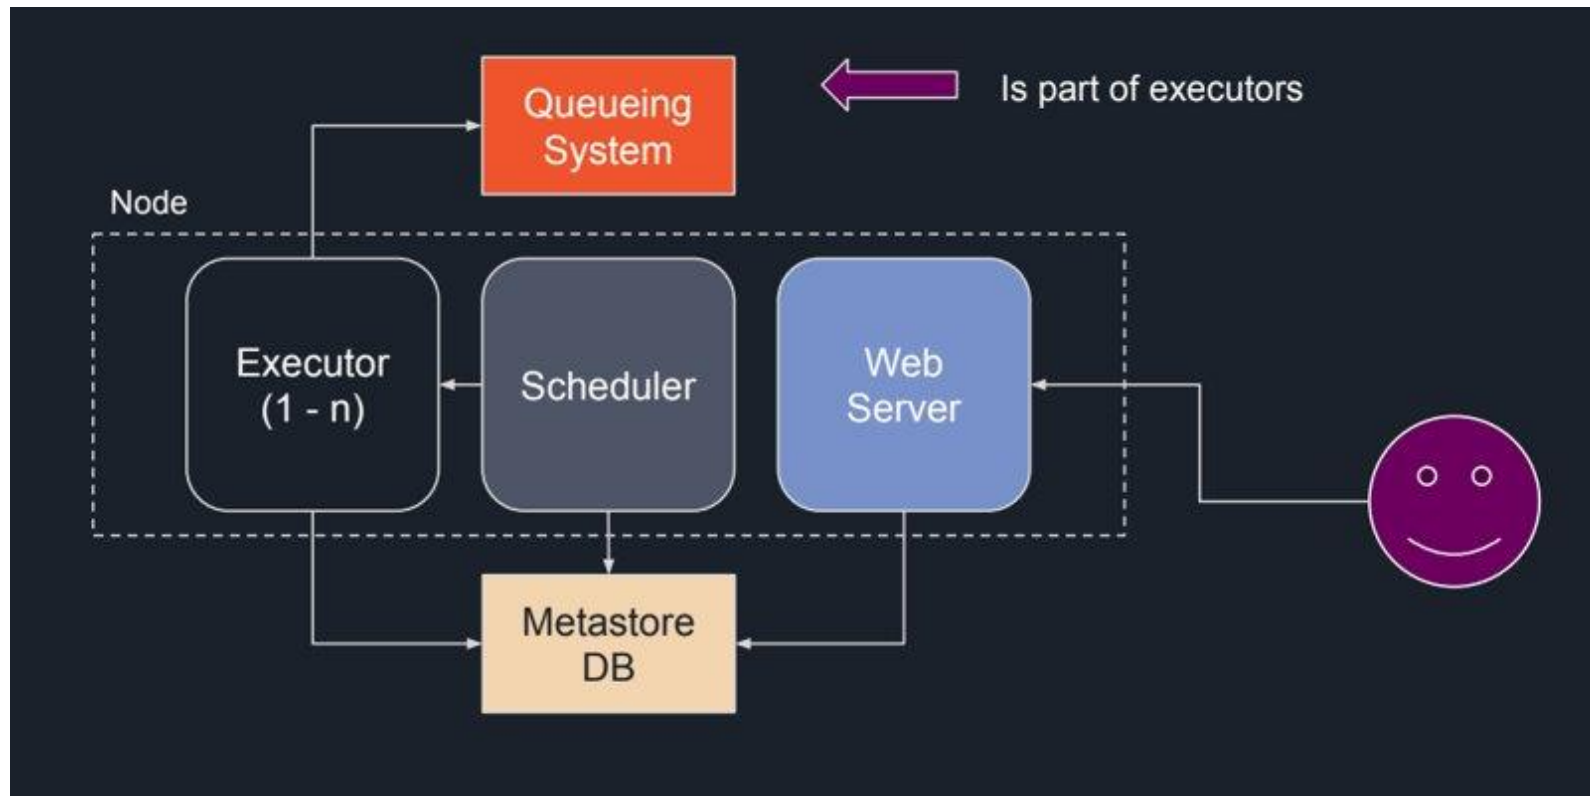
- Action Operators
- Transfer Operators
- Sensor Operators

- TASK
- TASK INSTANCE
- DEPENDENCIES

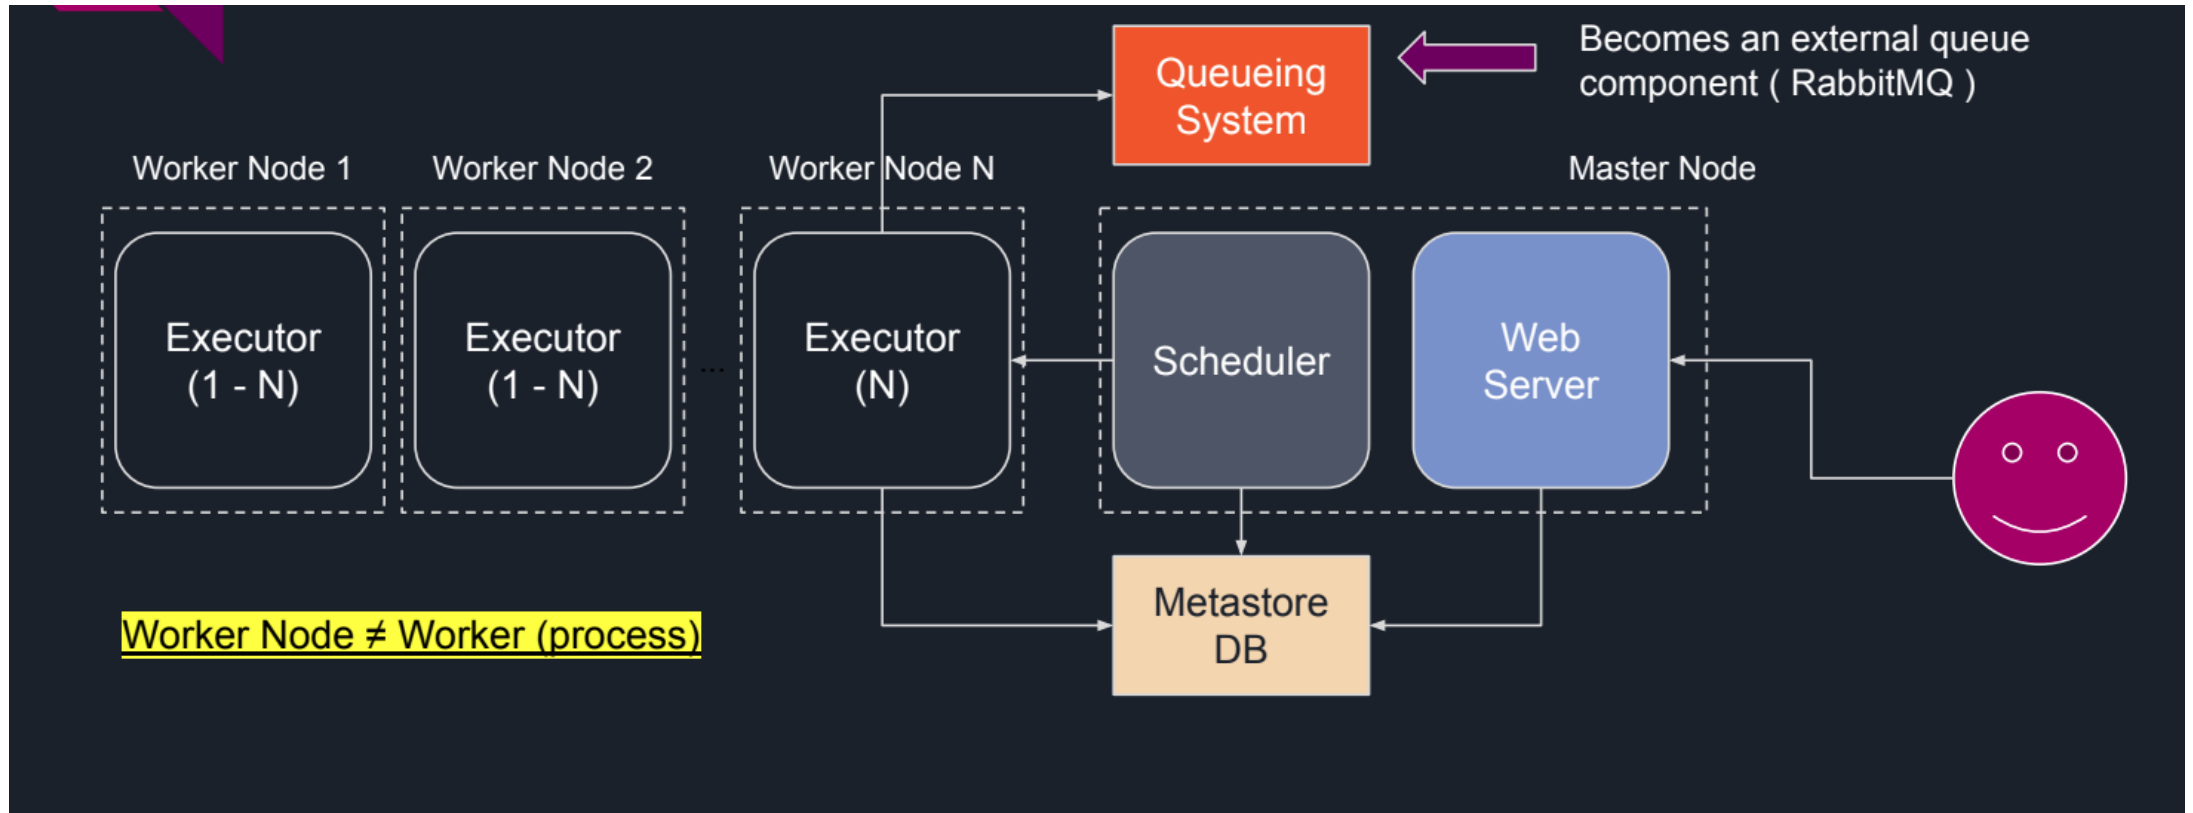


- WORKFLOW

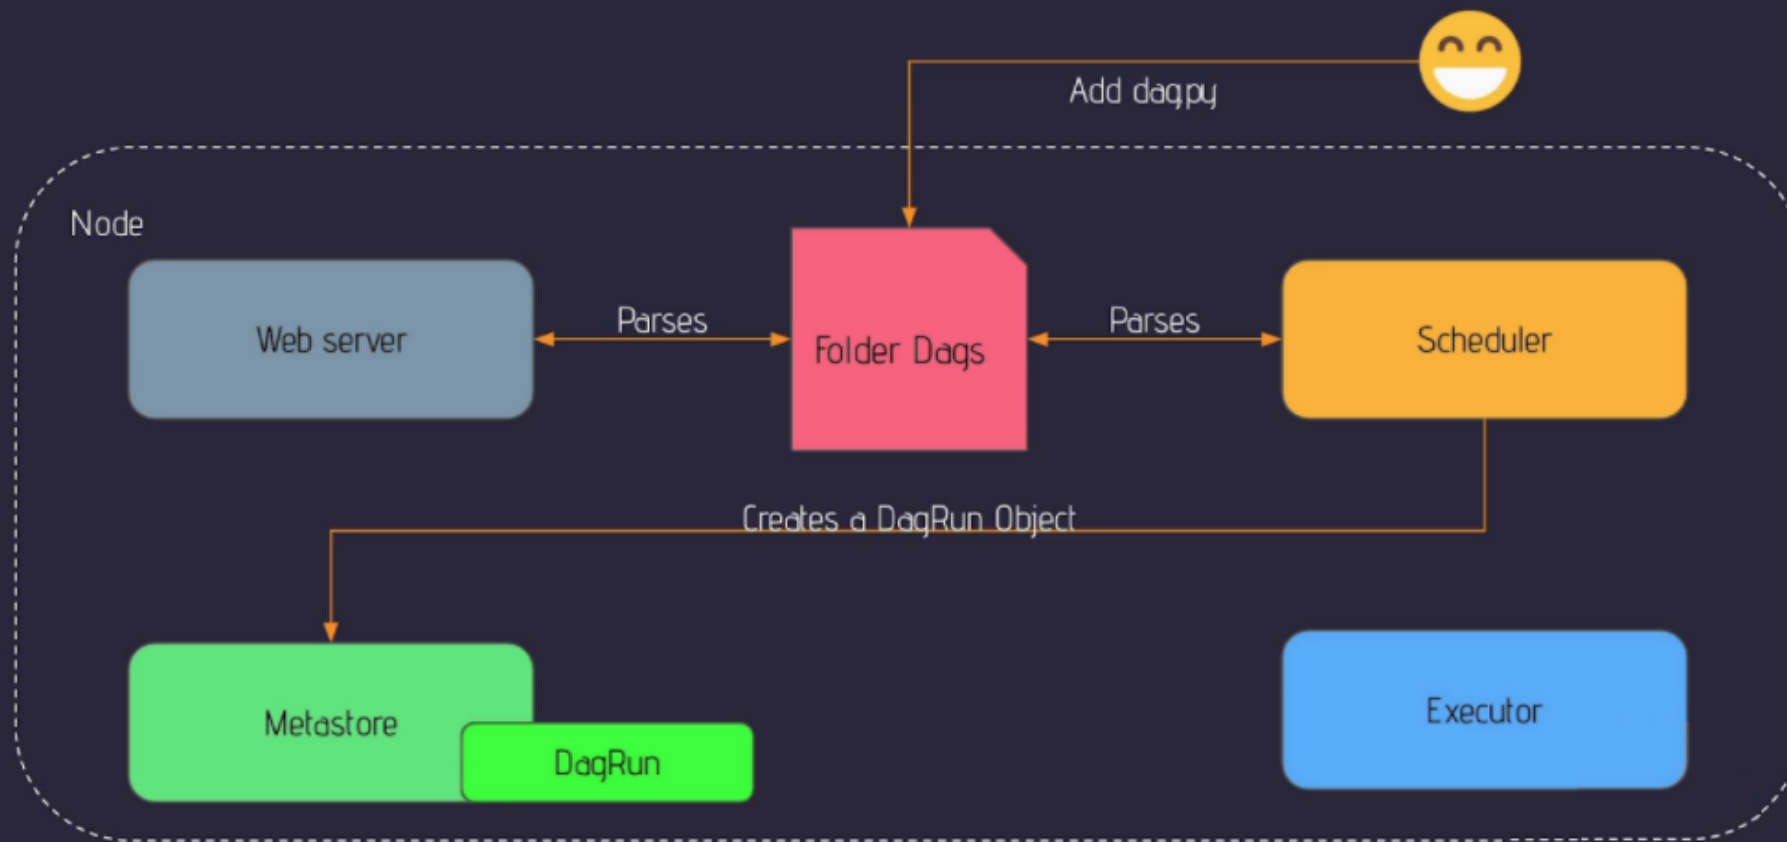




Airflow
Single Node
Architecture

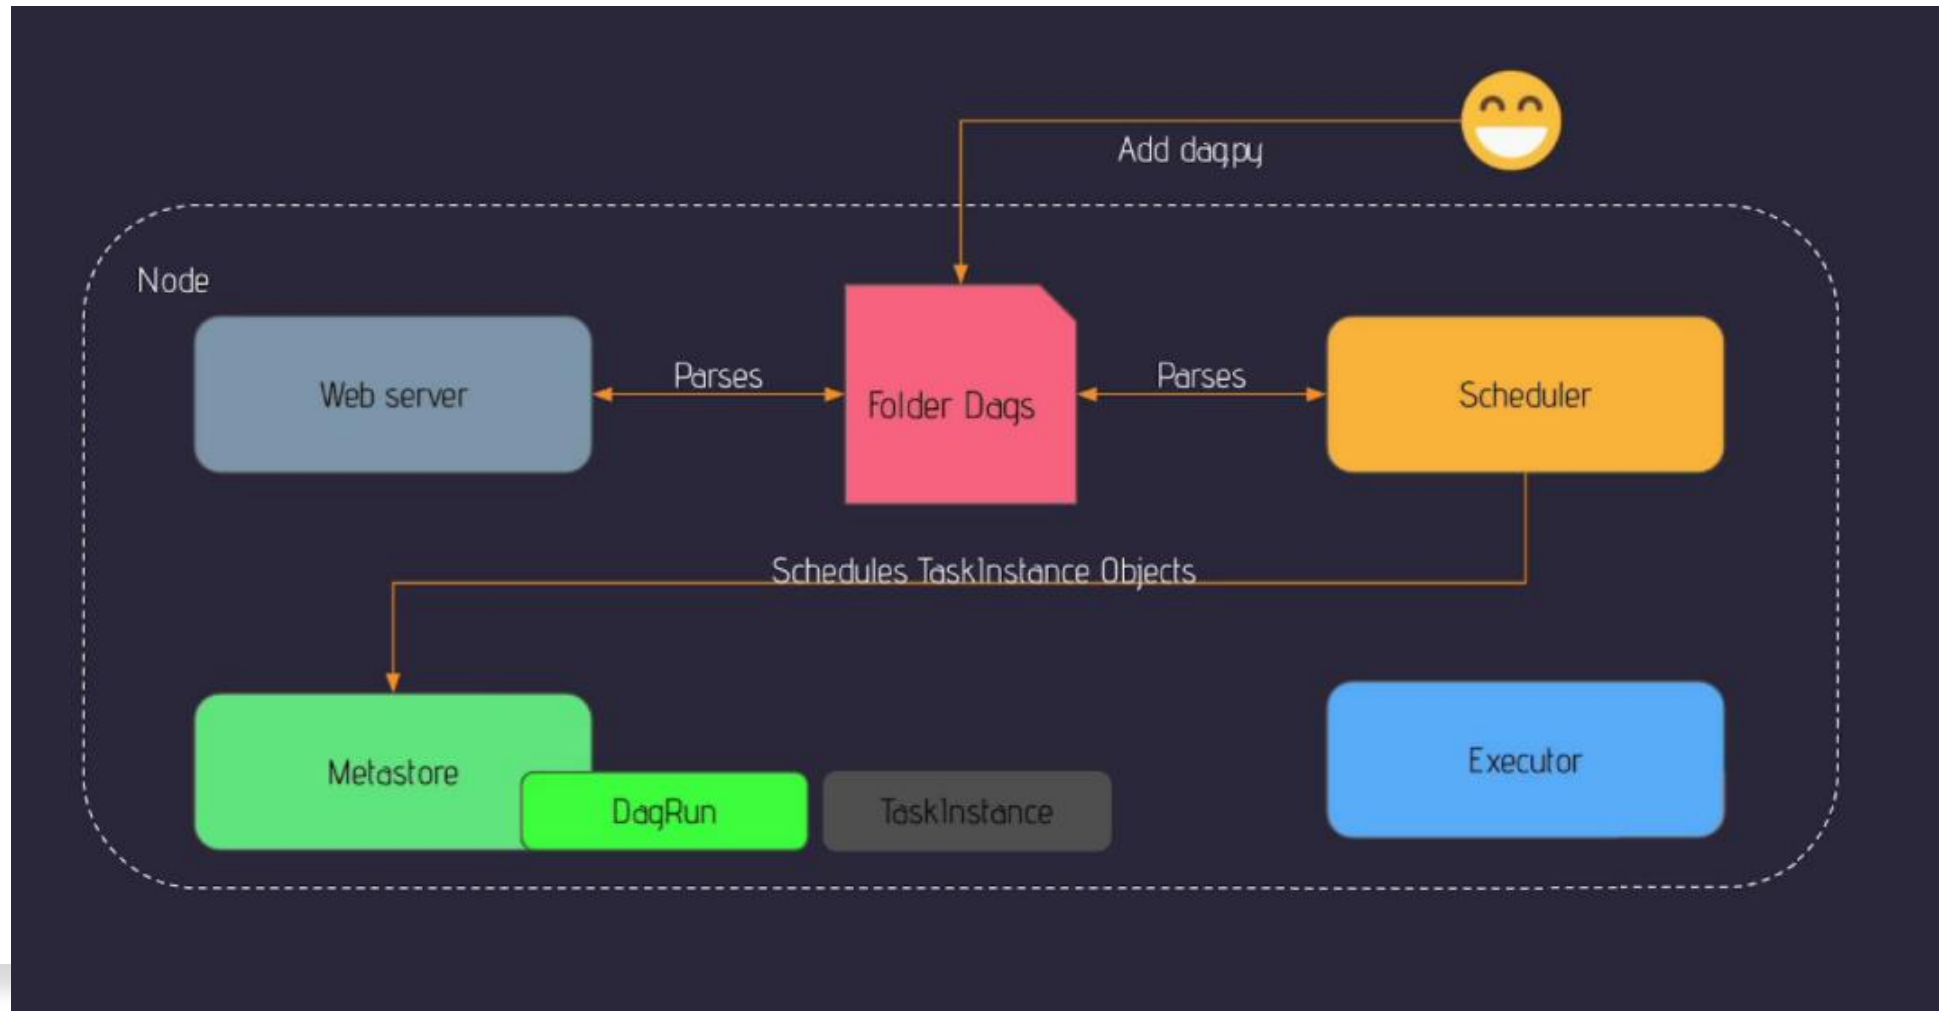


Airflow Multinode Architecture



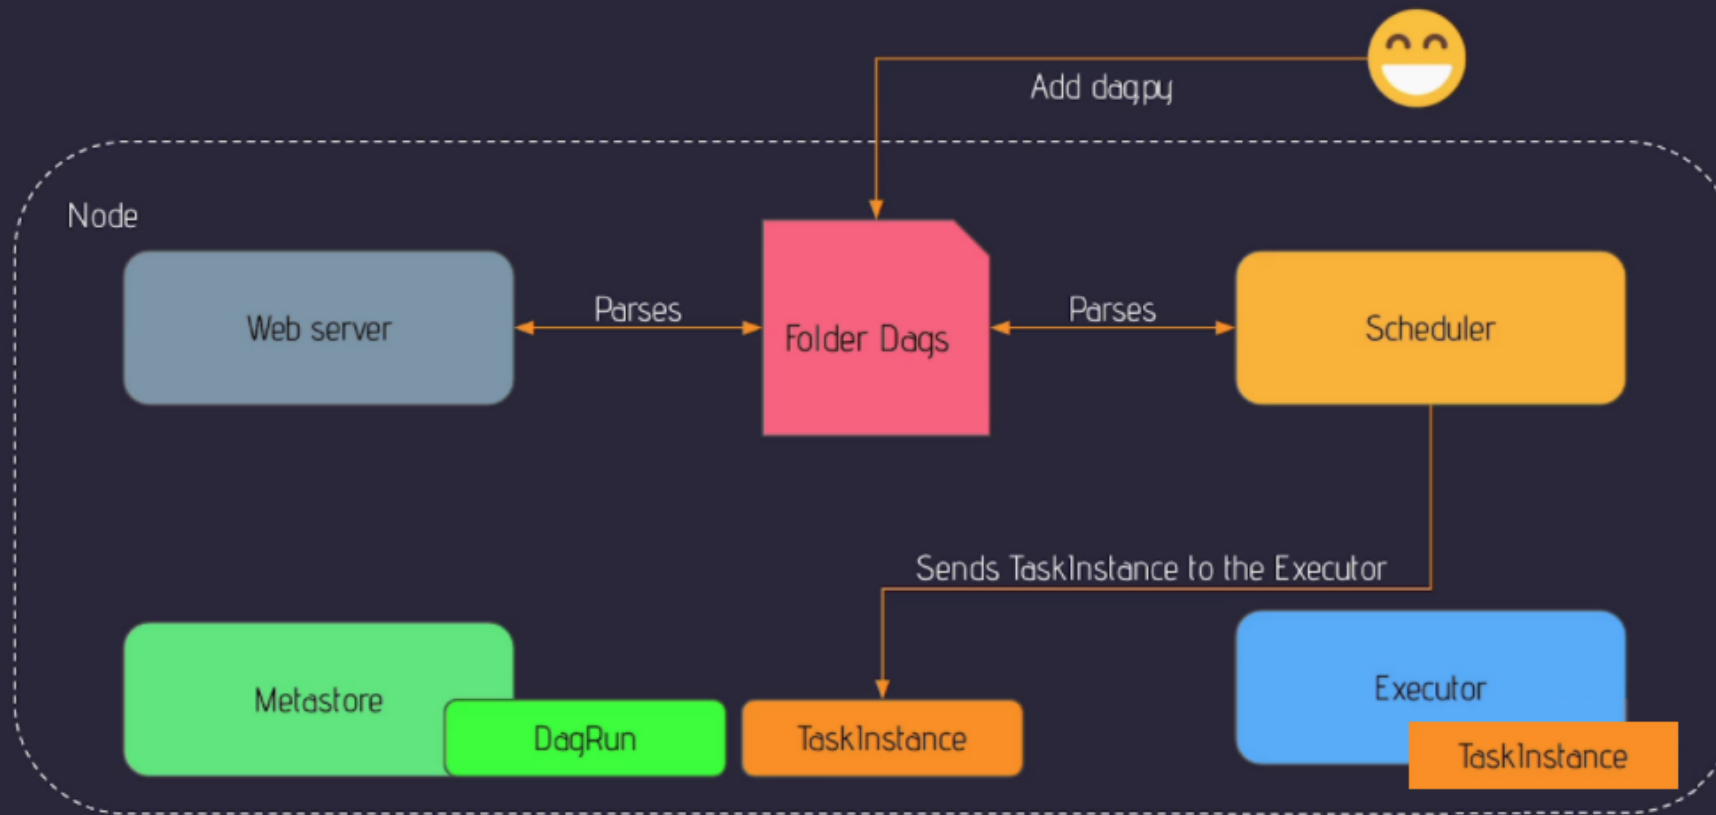
How Airflow Works?

- Upload
- DAG Parsing – webserver and scheduler
- DagRun Object in Metadata



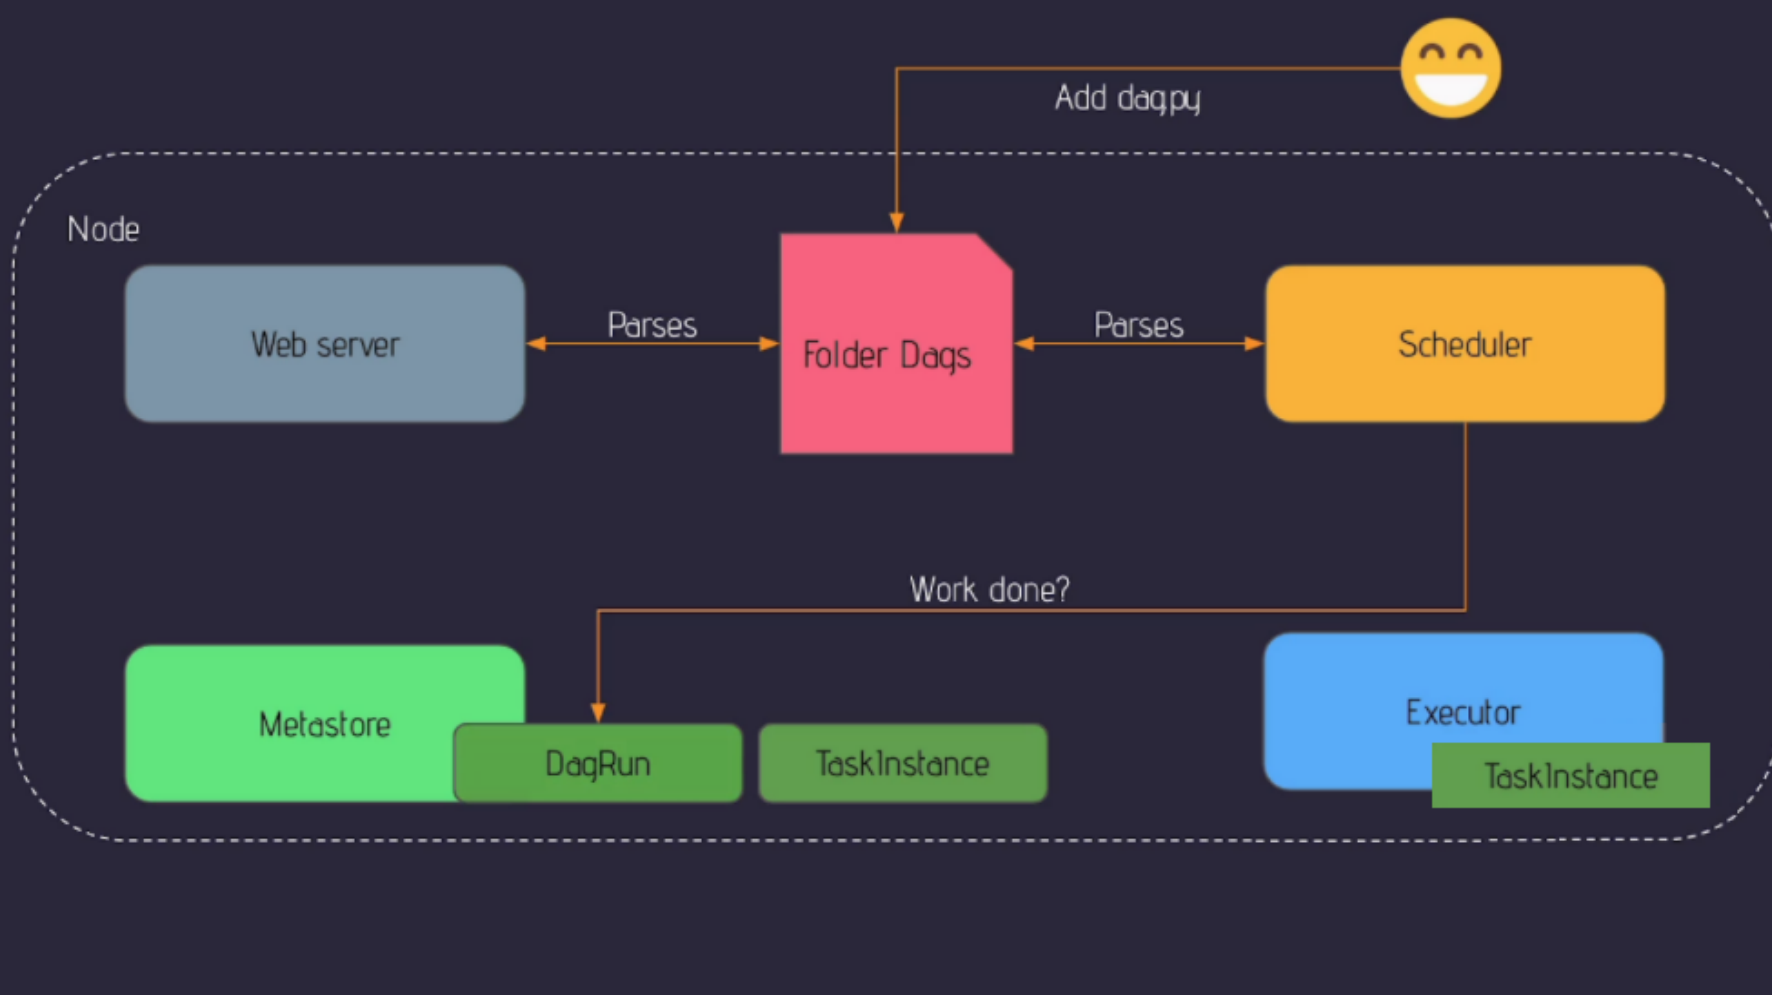
How Airflow Works?

- Task Instance is created by Scheduler,
 - No Status in the beginning.
 - Scheduled, when ready to be triggered



How Airflow Works?

- **Task Instance** status will be **Queued at this time**.
- Queued tasks are executed into the worker. Task state **running**.
- Executor updates the Task instance, status **running**, **failed** or **success** in the metastore.



How Airflow Works?

- Scheduler checks if more tasks exist? If all completed, DagRun is marked as success/failed
- In parallel to all of this, Web UI is being updated about the status of the task.
- **Task Lifecycle.**

No status → scheduled → queued → running → success/failed.

- Quick Tour of Web UI
- Quick Tour of Airflow CLI
- DAG and Operator definition in python
- Dependencies between tasks
- Connections
- Variables
- Extras and Provider Packages (Operators and hooks)
- XComs
- Branching
- Trigger Rules
- Sub Dags
- Taskgroups
- Task flow API

Quick Tour of Web UI localhost:8080

Quick Tour of Airflow CLI

NOTE : airflow 2.X has updated the CLI usage.

- airflow db init
- airflow db upgrade
- airflow db reset
- airflow webserver
- airflow scheduler
- airflow dags unpause <dag_id>
- airflow dags pause <dag_id>
- airflow dags trigger -e <execution_date> (-e optional)
- airflow dags list
- airflow tasks list <dag_id>
- airflow tasks test <dag_id> <task_id> <execution_date>
- airflow dags backfill -s <start_date> -e <end_date> --reset_dag_runs <dag_id>

DAG and Operator definition in python

Dependencies between tasks

- $t1 \gg t2$
- $t2 \ll t1$
- `t2.set_upstream(t1)`
- `t1.set_downstream(t2)`
- Chain dependency
- Cross dependency

Airflow Connections

- FROM UI
- FROM REST API
- FROM CLI
 - `airflow connections add 'my_prod_db' \`
 - `--conn-type 'my-conn-type'`
 - `...`
 -
- Export Connections from CLI
 - `airflow connections export connections.json`
 - `airflow connections export /tmp/connections --format yaml`
- Store Connection in Environment Variables
 - `.bashrc`
 - `Docker.env`

Airflow Variables

- How variables work?
- How to SET, GET a variable in Airflow?
- Optimizing variables with the JSON format
- Best practices with variables in Airflow

Airflow Variables

- How to hide the value of a variable?

```
DEFAULT_SENSITIVE_VARIABLE_FIELDS = (  
    'password',  
    'secret',  
    'passwd',  
    'authorization',  
    'api_key',  
    'apikey',  
    'access_token',  
)
```

PostgreSQL Connection and Operator

- Conn Type missing? Make sure you've installed the corresponding Airflow Provider Package.
- Provider Operator missing?
- Solution :
 - <https://airflow.apache.org/docs/apache-airflow/stable/extra-packages-ref.html>
 - `pip install 'apache-airflow[postgres]'`

Runtime Configs and XComs

- From airflow context
- Run time configs can be fetched from Dagrun object
- Xcoms -> Cross communication
- KEY VALUE TIMESTAMP
- Stored in Metadatabase, with an associated
 - Execution_date,
 - Task ID
 - DAG ID

Trigger Rules

- all_success
- all_failed
- all_done
- one_failed
- one_success
- none_failed
- none_failed_or_skipped – 1 parent succeeded
- none_skipped
- dummy